

Concurrency: Deadlock and Starvation Problems, more on Mutual Exclusion

Most slides have been adapted from «*Operating Systems: Internals and Design Principles*», 7/E W. Stallings (Chapter 6 + Appendix A).

Sistemi di Calcolo (Part II – Spring 2018)

Instructors: Daniele Cono D'Elia, Riccardo Lazzeretti

A real-world example

When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone. Statute passed by the Kansas State Legislature, early in the 20th century.



—A TREASURY OF RAILROAD FOLKLORE,
B. A. Botkin and Alvin F. Harlow


Deadlock

- The permanent blocking of a set of processes that either compete for system resources or communicate with each other
- A set of processes is deadlocked when each process in the set is blocked awaiting an event that can only be triggered by another blocked process in the set
- Permanent
- No efficient solution



Potential Deadlock

I need quad
C and D




c

I need quad
B and C




b

I need quad
D and A



d

I need quad
A and B



a


Actual Deadlock

HALT until D
is free

HALT until C
is free

HALT until B
is free

HALT until A
is free



Resource Categories

Reusable resource

Can be safely used by only one process at a time, and it is not depleted by that use

Examples: processors, I/O channels, main and secondary memory, devices, and structures such as files, databases, and semaphores

Consumable resource

One that can be created (produced) and destroyed (consumed)

Examples: interrupts, signals, messages, and information in I/O buffers

Reusable Resources - Example

Process P

Step	Action
p ₀	Request (D)
p ₁	Lock (D)
p ₂	Request (T)
p ₃	Lock (T)
p ₄	Perform function
p ₅	Unlock (D)
p ₆	Unlock (T)

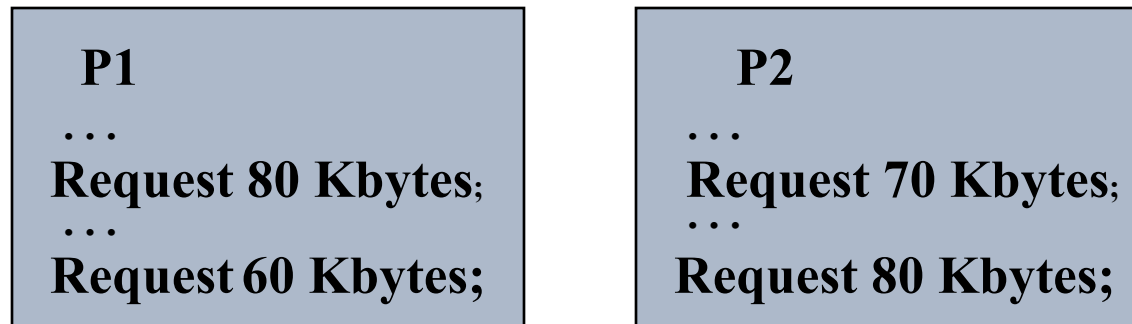
Process Q

Step	Action
q ₀	Request (T)
q ₁	Lock (T)
q ₂	Request (D)
q ₃	Lock (D)
q ₄	Perform function
q ₅	Unlock (T)
q ₆	Unlock (D)

Figure 6.4 Example of Two Processes Competing for Reusable Resources

Example 2: Memory Request


- Space is available for allocation of 200Kbytes, and the following sequence of events occur:



- Deadlock occurs if both processes progress to their second request

Consumable Resources Deadlock

- Consider a pair of processes, in which each process attempts to receive a message from the other process and then send a message to the other process:



- Deadlock occurs if the Receive is blocking

Conditions for Deadlock

Mutual Exclusion	Hold-and-Wait	No Pre-emption	Circular Wait
<ul style="list-style-type: none">• only one process may use a resource at a time	<ul style="list-style-type: none">• a process may hold allocated resources while awaiting assignment of others	<ul style="list-style-type: none">• no resource can be forcibly removed from a process holding it	<ul style="list-style-type: none">• a closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain

The first three conditions are necessary but not sufficient

Dealing with Deadlock

- Three general approaches exist for dealing with deadlock:

Prevent Deadlock

- adopt a policy that eliminates one of the conditions

Avoid Deadlock

- make the appropriate dynamic choices based on the current state of resource allocation

Detect Deadlock

- attempt to detect the presence of deadlock and take action to recover

Deadlock Prevention Strategy

- Design a system in such a way that the possibility of deadlock is excluded
- Two main methods:
 - Indirect
 - prevent the occurrence of one of the three necessary conditions
 - Direct
 - prevent the occurrence of a circular wait

Deadlock Condition Prevention

Mutual Exclusion

if access to a resource requires mutual exclusion then it must be supported by the OS

Hold and Wait

require that a process request all of its required resources at one time and blocking the process until all requests can be granted simultaneously

Deadlock Condition Prevention


- **No Preemption**

- if a process holding certain resources is denied a further request, that process must release its original resources and request them again; alternatively, OS may preempt the second process and require it to release its resources
- approach practical only for resources for which the state can be saved and later restored (e.g., CPU)

- **Circular Wait**

- define a linear ordering of resource types
- if a resource of type R has been assigned to a process, the latter can only request resources whose types follow R in the order

Deadlock Avoidance



- A decision is made dynamically whether the current resource allocation request will, if granted, potentially lead to a deadlock
- Requires knowledge of future process requests
- Allows more concurrency than deadlock prevention



Two Approaches to Deadlock Avoidance


Deadlock Avoidance

Resource Allocation Denial

- do not grant an incremental resource request to a process if this allocation might lead to deadlock


Process Initiation Denial

- do not start a process if its demands might lead to deadlock



Deadlock Avoidance Advantages

- Process initiation denial is hardly optimal...
- ...while resource allocation denial has a few perks!
 - It is not necessary to preempt and roll back processes, as in deadlock detection (*we will discuss detection soon*)
 - It is less restrictive than deadlock prevention



Deadlock Avoidance Restrictions

- Maximum resource requirement for each process must be stated in advance
- Processes under consideration must be independent and with no synchronization requirements
- There must be a fixed number of resources to allocate
- No process may exit while holding resources

Deadlock Strategies

Deadlock prevention strategies are very conservative

- limit access to resources by imposing restrictions on processes

Deadlock detection strategies do the opposite

- resource requests are granted whenever possible

Deadline Detection Algorithms

A check for deadlock can be made as frequently as each resource request or, less frequently, depending on how likely it is for a deadlock to occur.

Detection performed at each resource request:

- Pros: leads to early detection, algorithm is simple
- Cons: frequent checks consume considerable CPU time

Deadline Detection - Recovery


Some approaches, in increasing order of sophistication:

- Abort all deadlocked processes (most common)
- Back up each deadlocked process to some previously defined checkpoint and restart all processes
- Successively abort deadlocked processes until deadlock no longer exists
- Successively preempt resources until deadlock no longer exists (requires rollback mechanism)

Dining Philosophers Problem

No two philosophers can use the same fork at the same time (mutual exclusion)

No philosopher must starve (avoid deadlock & starvation)



Why so relevant?

- illustrates basic problems in deadlock and starvation
- standard test case for evaluating approaches to synchronization

Using N semaphores?

```
semaphore fork [5] = {1};  
void philosopher (int i)  
{  
    while (true) {  
        think();  
        wait (fork[i]);  
        wait (fork [(i+1) mod 5]);  
        eat();  
        signal(fork [(i+1) mod 5]);  
        signal(fork[i]);  
    }  
}
```

Try to pick the i -th fork on the left, then $(i+1)$ -th fork on the right: the code above will easily lead to deadlock

A Correct Solution . . .

```
semaphore fork[5] = {1};  
semaphore room = {4};  
void philosopher (int i)  
{  
    while (true) {  
        think();  
        wait (room);  
        wait (fork[i]);  
        wait (fork [(i+1) mod 5]);  
        eat();  
        signal (fork [(i+1) mod 5]);  
        signal (fork[i]);  
        signal (room);  
    }  
}
```


Software approaches to mutual exclusion

Mutual exclusion can be implemented at software level too

Scenario: processes communicate via a central memory on a single/multi-processor machine

Assumption: mutual exclusion at the memory access level

- read/write operations for the same location are serialized by some sort of memory arbiter (the order is unknown)

First attempt

```
/* global */  
int turn = 0;  
  
// assignments valid for P0 (flip for P1)  
int me = 0, other = 1;  
  
while (turn != me)  
    /* busy wait */ ;  
/* CS */  
turn = other;
```

Second attempt

```
/* global */  
boolean flag[2] = {false, false};  
  
// assignments valid for P0 (flip for P1)  
int me = 0, other = 1;  
  
while (flag[other])  
    /* busy wait */ ;  
flag[me] = true;  
/* CS */  
flag[me] = false;
```

Third attempt

```
// assignments valid for P0 (flip for P1)
int me = 0, other = 1;

flag[me] = true;
while (flag[other])
    /* busy wait */ ;
/* CS */
flag[me] = false;
```

Fourth attempt

```
// assignments valid for P0 (flip for P1)
int me = 0, other = 1;

flag[me] = true;
while (flag[other]) {
    flag[me] = false;
    /* delay */
    flag[me] = true;
}
/* CS */
flag[me] = false;
```

Correct Solution: Dekker's Algorithm

```
int me = 0, other = 1; // P0 (flip for P1)

while (true) {
    flag[me] = true;
    while (flag[other]) {
        if (turn == other) {
            flag[me] = false;
            while (turn == other) /* busy wait */ ;
            flag[me] = true;
        }
    }
    /* CS */
    turn = other;
    flag[me] = false;
}
```

Dekker's Algorithm, informally

```
flag[me] = true;
while (flag[other]) {
    if (turn==other) {
        flag[me] = false;
        while (turn==other);
        flag[me] = true;
    }
}
/* CS */
turn = other;
flag[me] = false;
```

*I want to enter
If you want to enter
and if it's your turn
I don't want any more
If it's your turn I'll wait
I want to enter*

*You can enter next
I don't want any more*

Credits: <https://cs.stackexchange.com/q/12621>

Dijkstra (1965)

Dijkstra generalized Dekker's algorithm to deal with N entities

He also provided definitions for mutual exclusion (ME), no deadlock (ND), and no starvation (NS) properties

=> notice that NS implies ND

```
/* global storage */
boolean interested[N] = {false, ..., false}
boolean passed[N]      = {false, ..., false}
int k = <any>           // k ∈ {0, 1, ..., N-1}

/* local info */
int i = <entity ID>     // i ∈ {0, 1, ..., N-1}
```


Dijkstra's Algorithm

```
1. interested[i] = true
2. while (k != i) {
3.     passed[i] = false
4.     if (!interested[k]) then k = i
5. }
5. passed[i] = true
6. for j in 1 ... N except i do
7.     if (passed[j]) then goto 2
8. <critical section>
9. passed[i] = false; interested[i] = false
```

A closer look

Trying protocol

```
1. interested[i] = true
2. while (k != i) {
3.     passed[i] = false
4.     if (!interested[k]) then k = i
5. }
5. passed[i] = true
6. for j in 1 ... N except i do
7.     if (passed[j]) then goto 2
```

1: process i shows interest in entering CS

2-4: k “chooses” among processes that showed interest

5: phase one passed by process i

6-7: restart if more than one process passed phase one

Example ($N \geq 4$, initially $k = 4$)

P1	P2	P3
<code>intr[1] = true</code>		
<code>while (k!=1)</code>		
<code>pass[1] = false</code>		
	<code>intr[2] = true</code>	
	<code>while (k!=2)</code>	
	<code>pass[2] = false</code>	
		<code>intr[3] = true</code>
		<code>while (k!=3)</code>
		<code>pass[3] = false</code>
<code>if (!intr[4])</code>		
	<code>if (!intr[4])</code>	
		<code>if (!intr[4])</code>
<code>k = 1</code>		
	<code>k = 2</code>	
		<code>k = 3</code>

Entering the Critical Section

Processes P1, P2, and P3 are now ready to execute lines 5-7

- only one of them will be able to access CS
- **k=3** does not imply that P3 goes first: k is used only to solve conflicts, thus the scheduler may also let P1 or P2 in first!
- a conflict arises when two processes set `pass[i]=true` before the other has entered and then left the critical section

P3
<code>pass[3] = true</code>
<code>for j in 1..N except 3 do</code>
<code>if pass[j] goto 2</code> <i>(skipped)</i>
<code><critical section></code>
<code>pass[3] = false</code>
<code>intr[3] = false</code>

One possible conflict-free continuation: the scheduler lets P3 run with P1 and P2 still “stopped” at line 5

P3 finished, conflict for P1 & P2?

P1	P2
while (k!=1)	
pass[1] = false	
	intr[2] = true
	while (k!=2)
if (!intr[3]) <i>(true: P3 done)</i>	
	if (!intr[3]) <i>(true: P3 done)</i>
then k = 1	
while (k!=1) <i>(skipped)</i>	
pass[1] = true	
	then k = 2
	while (k!=2) <i>(skipped)</i>
if (pass[2]) goto 2 <i>(taken!!)</i>	
	if (pass[1]) goto 2 <i>(taken!!)</i>

When a conflict occurs, the algorithm goes back to line 2 and clears pass[i]... except for the process that set k last! (P2 here)